



# Information Retrieval

Amin Nazari

Spring 2025

# Chapter 8

Modern text classification

# Agenda

- Introduction to Text Classification
- MLP for Text Classification
- RNN for Text Classification
- LSTM for Text Classification
- Comparison & Results
- Conclusion

# What is Text Classification?

- Assigning categories to text data
- Applications: Spam detection, sentiment analysis, topic labeling
- Challenges: High-dimensional data, sequence dependency, contextual meaning

# Preprocessing for Text Classification

- Tokenization
- Lowercasing
- Stopword Removal
- Vectorization (TF-IDF, Word2Vec, Embeddings)


# Preprocessing for Text Classification

	Document 1	Document 2	Document 3	Document 4	Document 5	Document 6	Document 7	Document 8
Term(s) 1	10	0	1	0	0	0	0	2
Term(s) 2	0	2	0	0	0	18	0	2
Term(s) 3	0	0	0	0	0	0	0	2
Term(s) 4	6	0	0	4	6	0	0	0
Term(s) 5	0	0	0	0	0	0	0	2
Term(s) 6	0	0	1	0	0	1	0	0
Term(s) 7	0	1	8	0	0	0	0	0
Term(s) 8	0	0	0	0	0	3	0	0



← Word Vector  
(Passage Vector)

Document Vector

# Preprocessing for Text Classification




# Basic Components of Perceptron






# XOR


AND		
$I_1$	$I_2$	out
0	0	0
0	1	0
1	0	0
1	1	1




OR		
$I_1$	$I_2$	out
0	0	0
0	1	1
1	0	1
1	1	1



XOR		
$I_1$	$I_2$	out
0	0	0
0	1	1
1	0	1
1	1	0



# Forward Propagation



here,

$x$  - inputs

$w_{ij}^n$  -  $n$  is num of next hidden layer,

$i$  is neuron num of previous layer,

$j$  is neuron num of next layer

( $w_{21}^2$  -  $n$  is next layer,

$i$  is 2<sup>nd</sup> neuron of previous layer,

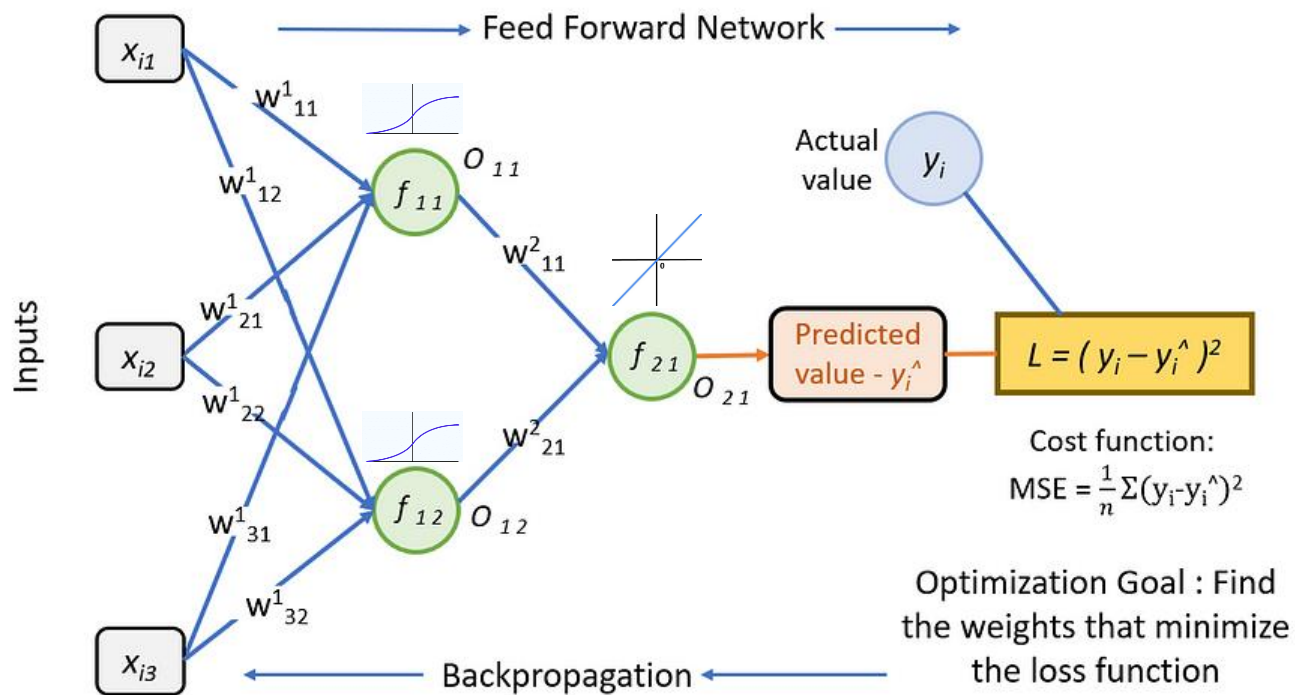
$j$  is 1<sup>st</sup> neuron of next layer)

$f_{ij}$  -  $i$  is num of layer and  $j$  is num of neuron

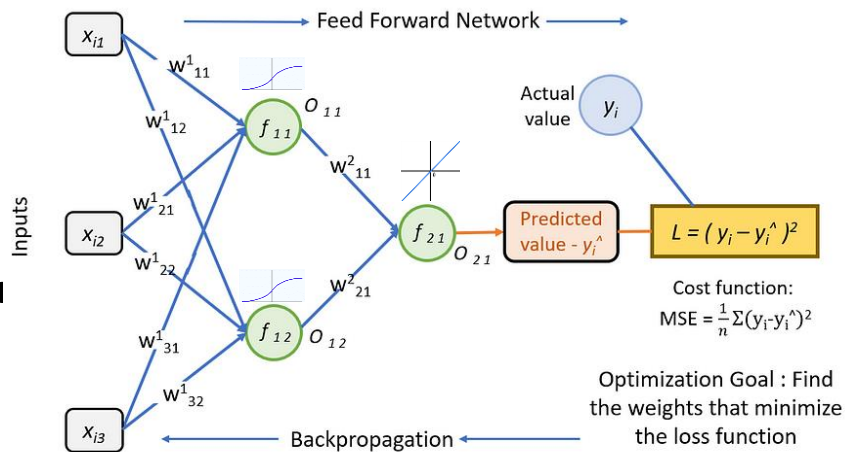
$O_{ij}$  - output from  $i^{\text{th}}$  layer  $j^{\text{th}}$  neuron

$Y$  - output

# Back Propagation



# Back Propagation (Regression)



$$f_{11} = w_{11} * x_1 + w_{21} * x_2 + w_{31} * x_3$$

$$f_{12} = w_{12} * x_1 + w_{22} * x_2 + w_{32} * x_3$$

$$O_{11} = \text{Sigmoid}(f_{11})$$

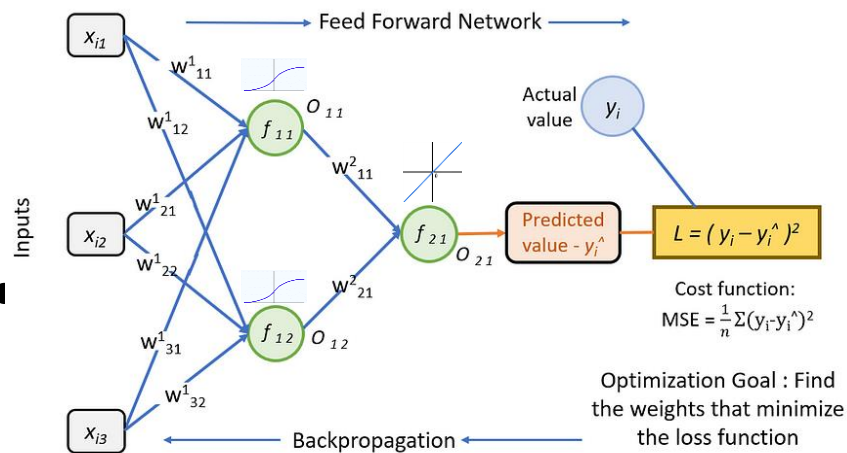
$$O_{12} = \text{Sigmoid}(f_{12})$$

$$f_{21} = v_{11} * O_{11} + v_{21} * O_{21}$$

$$O_{21} = f_{21}$$

$$E = (y - O_{21})^2$$

# Back Propagation (Regression)



$$f_1 = x_{(1*3)} * w_{(3*2)}$$

$$O_1 = \text{Sigmoid}(f_1)$$

$$f_2 = O_{1(1*2)} * v_{(2*1)}$$

$$O_2 = (f_2)$$

Cost function:  
 $MSE = \frac{1}{n} \sum (y_i - y_i^{\wedge})^2$

Optimization Goal : Find  
 the weights that minimize  
 the loss function

$$E = (y - O_2)^2$$

# Back Propagation (Regression)


- $w = w - \alpha \frac{\partial E(w)}{\partial w}$
- $v = v - \alpha \frac{\partial E(v)}{\partial v}$
- $f_1 = x_{(1 \times 3)} * w_{(3 \times 2)}; O_1 = \text{Sigmoid}(f_1)$
- $f_2 = O_{1(1 \times 2)} * v_{(2 \times 1)}; O_2 = f_2;$
- $E = (y - O_2)^2;$
- $\frac{\partial E}{\partial v} = \frac{\partial E}{\partial O_2} \cdot \frac{\partial O_2}{\partial f_2} \cdot \frac{\partial f_2}{\partial v} = -(y - O_2) \times 1 \times O_1$
- $\frac{\partial E}{\partial w} = \frac{\partial E}{\partial O_2} \cdot \frac{\partial O_2}{\partial f_2} \cdot \frac{\partial f_2}{\partial o_1} \cdot \frac{\partial o_1}{\partial f_1} \cdot \frac{\partial f_1}{\partial w} = -(y - O_2) \times 1 \times v \times (O_1(1 - O_1)) \times x$

# Back Propagation (Classification)

- $w = w - \alpha \frac{\partial E(w)}{\partial w}$
- $v = v - \alpha \frac{\partial E(v)}{\partial v}$
- $f_1 = x_{(1*3)} * w_{(3*2)}; O_1 = \text{Sigmoid}(f_1)$
- $f_2 = O_{1(1*2)} * v_{(2*1)}; O_2 = \text{Sigmoid}(f_2);$
- $E = y * \log(O_2) + (1 - y) * \log(1 - O_2);$
- $\frac{\partial E}{\partial v} = \frac{\partial E}{\partial O_2} \cdot \frac{\partial O_2}{\partial f_2} \cdot \frac{\partial f_2}{\partial v} = -(\frac{y}{O_2} + \frac{-(1-y)}{1-O_2}) \times (O_2(1-O_2)) \times O_1 = -(y - O_2) \times O_1$
- $\frac{\partial E}{\partial w} = \frac{\partial E}{\partial O_2} \cdot \frac{\partial O_2}{\partial f_2} \cdot \frac{\partial f_2}{\partial o_1} \cdot \frac{\partial o_1}{\partial f_1} \cdot \frac{\partial f_1}{\partial w} = -(\frac{y}{O_2} + \frac{-(1-y)}{1-O_2}) \times (O_2(1-O_2)) \times v \times (O_1(1-O_1)) \times x$
- $\frac{\partial E}{\partial w} = -(y - O_2) \times v \times (O_1(1-O_1)) \times x$

# MLP (Multilayer Perceptron)

- Input  $\rightarrow$  Hidden Layers  $\rightarrow$  Output
- Uses fixed-size input (TF-IDF, bag-of-words)
- Limitation: Ignores word order and sequence

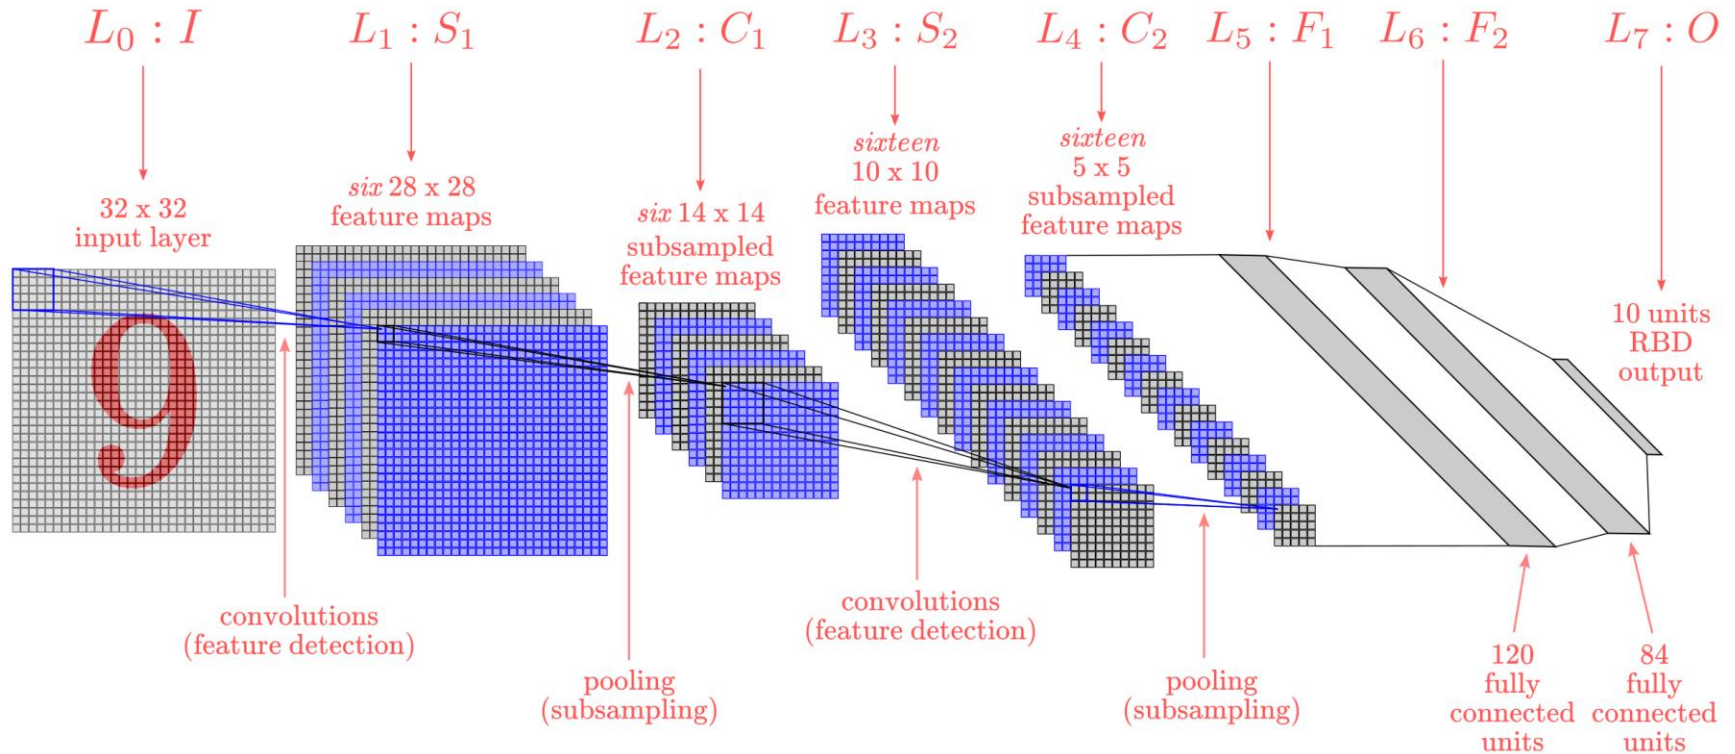




# Convolutional Neural Networks (CNN)

- Originally designed for image recognition
- Captures local features and spatial hierarchies
- Effective for detecting patterns in text sequences
- Handles variable input length (via padding/truncation)

# Convolutional Neural Networks (CNN)



# CNN Architecture for Text

- Layers:
  - **Input:** Tokenized and embedded sentences (e.g., Word2Vec, GloVe)
  - **Embedding**
  - **Convolution Layer(s):** Filters slide over embedded text (like n-grams) and detects key patterns
  - **Pooling Layer(s):** Max pooling selects the most important features
  - **Fully Connected Layer**
  - **Softmax Output Layer**

# CNN Architecture for Text




Figure 3: General CNN architecture for text classification

# CNN Architecture for Text

- Input: Sentence of 100 words
- Embedding: 100 x 300 matrix
- Conv layer: 100 filters (kernel size = 3x300)
- Pooling: Global max pooling
- Dense: 128 neurons
- Output: Softmax with 3 classes


# CNN Architecture for Text

- To capture **different n-gram features**:
  - Filter size = 3 might capture short phrases like "free money now"
  - Filter size = 5 might capture longer phrases like "you have won a prize"

# RNN (Recurrent Neural Network)

- Designed for sequence data
- Maintains hidden state  $h_t$
- Processes text word by word
- Limitation: Vanishing gradient, poor long-term memory


# RNN (Recurrent Neural Network)





# RNN – Architecture

- RNN cells process sequence:  $x_1, x_2, \dots, x_n \rightarrow h_1, h_2, \dots, h_n$
- Output from final hidden state goes to classifier




# RNN – Architecture

For input sequence  $\{x_1, x_2, \dots, x_T\}$ , the hidden state is updated as:

$$h_t = \tanh(Wx_t + Uh_{t-1} + b)$$


Where:

- $W$ : Input weight matrix (same for all  $t$ )
- $U$ : Hidden-to-hidden weight matrix (also same)
- $b$ : Bias (shared)
- $h_t$ : Hidden state at time  $t$

 These matrices  $W, U, b$  are learned during training and reused at each time step.

- **Forward pass:** process sequence
- **Loss:** sum of losses at all time steps
- **Backward pass:** compute gradients at each time step and sum them up
- **Update:** use the summed gradients to update parameters once per sequence (or batch)

$$\frac{\partial L}{\partial \theta} = \sum_{t=1}^T \frac{\partial L_t}{\partial \theta}$$



# Why Do We Need LSTM?

- Traditional RNNs struggle with long sequences.
- They forget important past information due to:
  - Vanishing gradients
  - Exploding gradients
- Example: Predicting words in long sentences → needs memory of earlier context.


# LSTM in a Nutshell

- A type of Recurrent Neural Network (RNN).
- Designed to remember long-term dependencies.
- Introduced by Hochreiter & Schmidhuber (1997).
- Uses gates to control memory.

# LSTM Cell: Overview

- Maintains two states:
  - Hidden state ( $h_t$ )
  - Cell state ( $c_t$ )
- Controls memory using three gates:
  - Forget gate
  - Input gate
  - Output gate

# LSTM



LSTM Model

# Step 1: Forget Gate

$$f_t = \sigma(W_f x_t + U_f h_{\{t-1\}} + b_f)$$

- Decides what information to discard from the cell state.
- Values close to 0  $\rightarrow$  forget
- Values close to 1  $\rightarrow$  keep



## Step 2: Input Gate

$$i_t = \sigma(W_i x_t + U_i h_{\{t-1\}} + b_i)$$
$$\tilde{c}_t = \tanh(W_c x_t + U_c h_{\{t-1\}} + b_c)$$

- Input gate: decides what new info to add.
- Candidate state: possible new values to add to memory.

## Step 3: Update Cell State

$$c_t = f_t \odot c_{\{t-1\}} + i_t \odot \tilde{c}_t$$

- Combine old memory with new info.
- Element-wise operations preserve gradient flow.

## Step 4: Output Gate

$$o_t = \sigma(W_o x_t + U_o h_{\{t-1\}} + b_o)$$
$$h_t = o_t \odot \tanh(c_t)$$

- Output gate: controls what to send to the next time step or output layer.

# Summary of LSTM Flow

- $x_t, h_{\{t-1\}}$   $\rightarrow$  all gates
- Gates control:
  - What to forget
  - What to update
  - What to output
- Cell state ( $c_t$ ) carries long-term memory
- Hidden state ( $h_t$ ) carries short-term output

# Why Use LSTM?

- Solves vanishing gradient problem
- Retains important info over long sequences
- Works well in:
  - Text generation
  - Language modeling
  - Time series prediction
  - Speech recognition

# Training Setup

- Models: KNN, SVM, ...
- Models: MLP (TF-IDF), CNN, RNN & LSTM

# Conclusion

- MLP: Simple, ignores word order
- RNN: Captures sequence, weak on long dependencies
- LSTM: Best performance via memory and context
- Future: Try BiLSTM, GRU, attention